



MySQL - RDBMS

Trainer: Mr. Nilesh Ghule




DBMS

- Any enterprise application need to manage data.
- In early days of software development, programmers store data into files and does operation on it. However data is highly application specific.
- Even today many software manage their data in custom formats e.g. Tally, Address book, etc.
- As data management became more common, DBMS systems were developed to handle the data. This enabled developers to focus on the business logic e.g. FoxPro, DBase, Excel, etc.
- At least CRUD (Create, Retrieve, Update and Delete) operations are supported by all databases.
Insert Select Edit Remov.
- Traditional databases are file based, less secure, single-user, non-distributed, manage less amount of data (MB), complicated relation management, file-locking and need number of lines of code to use in applications.



RDBMS

- RDBMS is relational DBMS.
- It organizes data into Tables, rows and columns. The tables are related to each other.
- RDBMS follow table structure, more secure, multi-user, server-client architecture, server side processing, clustering support, manage huge data (TB), built-in relational capabilities, table-locking or row-locking and can be easily integrated with applications.
- e.g. DB2, Oracle, MS-SQL, MySQL, MS-Access, SQLite, ...
- RDBMS design is based on Codd's rules developed at IBM (in 1970).



SQL - Structured Query Language (wiki)

- Clients send SQL queries to RDBMS server and operations are performed accordingly.
- Originally it was named as RQBE (Relational Query By Example).
- SQL is ANSI standardised in 1987 and then revised multiple times adding new features. Recent revision in 2016.
- SQL is case insensitive.
db name & table name
are case sensitive on Linux.
- There are five major categories:
 - DDL: Data Definition Language e.g. CREATE, ALTER, DROP, RENAME.
 - DML: Data Manipulation Language e.g. INSERT, UPDATE, DELETE. ← rows/records
 - DQL: Data Query Language e.g. SELECT. ← rows/records
 - DCL: Data Control Language e.g. CREATE USER, GRANT, REVOKE. ← security
 - TCL: Transaction Control Language e.g. SAVEPOINT, COMMIT, ROLLBACK. ← concurrent users
- Table & column names allows alphabets, digits & few special symbols.
- If name contains special symbols then it should be back-quotes.
- e.g. Tbl1, `T1#`, `T2\$` etc. Names can be ~~max~~ 30 chars long.



- Developed by Michael Widenius in 1995. It is named after his daughter name Myia.
- Sun Microsystems acquired MySQL in 2008.
- Oracle acquired Sun Microsystem in 2010.
- MySQL is free and open-source database under GPL. However some enterprise modules are close sourced and available only under commercial version of MySQL.
- MariaDB is completely open-source clone of MySQL.
- MySQL support multiple database storage and processing engines.
- MySQL versions:
 - < 5.5: MyISAM storage engine
 - 5.5: InnoDB storage engine
 - 5.6: SQL Query optimizer improved, memcached style NoSQL
 - 5.7: Windowing functions, JSON data type added for flexible schema
 - 8.0: CTE, NoSQL document store.
- MySQL is database of year 2019 (in database engine ranking).

MySQL → Community (Free)
Enterprise (Paid)

MySQL installation on Ubuntu/Linux

mysql-server

mysql-client

- terminal> sudo apt-get install mysql-community-server mysql-community-client


- This installs MySQL server (mysqld) and MySQL client (mysql).

- MySQL Server (mysqld)

- ✓ Run as background process.
- ✓ Implemented in C/C++.
- ✓ Process SQL queries and generate results.
- ✓ By default run on port 3306.
- Controlled via systemctl. (Ubuntu)
 - terminal> sudo systemctl start/stop/statusenable/disable mysql

(networking)

Socket
"ip addr
+ port"



- MySQL client (mysql)

- ✓ Command line interface
- ✓ Send SQL queries to server and display its results.
- terminal> mysql -u root -p

- Additional MySQL clients

- ✓ MySQL workbench
- ✓ PHPMyAdmin ✗



Getting started

- root login can be used to perform CRUD as well as admin operations.
- It is recommended to create users for performing non-admin tasks.
 - mysql> CREATE DATABASE db;
 - mysql> SHOW DATABASES;
 - mysql> CREATE USER dbuser@localhost IDENTIFIED BY 'dbpass';
 - mysql> SELECT user, host FROM mysql.user;
 - mysql> GRANT ALL PRIVILEGES ON db.* TO dbuser@localhost;
 - mysql> FLUSH PRIVILEGES;
 - mysql> EXIT;
- terminal> mysql –u dbuser –pdbpass
 - mysql> SHOW DATABASES;
 - mysql> SELECT USER(), DATABASE();
 - mysql> USE db;
 - mysql> SHOW TABLES;
 - mysql> CREATE TABLE student(id INT, name VARCHAR(20), marks DOUBLE);
 - mysql> INSERT INTO student VALUES(1, 'Abc', 89.5);
 - mysql> SELECT * FROM student;



Database physical layout

- In MySQL, the data is stored on disk in its data directory i.e. /var/lib/mysql
- Each database/schema is a separate sub-directory in data dir.
- Each table in the db, is a file on disk.
- e.g. student table in current db is stored in file /var/lib/mysql/db/student.ibd.
- Data is stored in binary format.
- A file may not be contiguously stored on hard disk.
- Data rows are not contiguous. They are scattered in the hard disk.
- In one row, all fields are consecutive.
- When records are selected, they are selected in any order.



INSERT – DML

- Insert a new row (all columns, fixed order).
 - `INSERT INTO table VALUES (v1, v2, v3);`
- Insert a new row (specific columns, arbitrary order).
 - `INSERT INTO table(c3, c1, c2) VALUES (v3, v1, v2);`
 - `INSERT INTO table(c1, c2) VALUES (v1, v2);`
 - Missing columns data is `NULL`.
 - `NULL` is special value and it is not stored in database.
- Insert multiple rows.
 - `INSERT INTO table VALUES (av1, av2, av3), (bv1, bv2, bv3), (cv1, cv2, cv3).`
- Insert rows from another table.
 - `INSERT INTO table SELECT c1, c2, c3 FROM another-table;`
 - `INSERT INTO table (c1,c2) SELECT c1, c2 FROM another-table;`





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>






MySQL - RDBMS

Trainer: Mr. Nilesh Ghule



Database logical layout


- Database/schema is like a namespace/container that stores all db objects related to a project.
- It contains tables, constraints, relations, stored procedures, functions, triggers, ...
- There are some system databases e.g. mysql, performance_schema, information_schema, sys, ... They contains db internal/system information.
 - e.g. SELECT user, host FROM mysql.user;
- A database contains one or more tables.
- Tables have multiple columns.
- Each column is associated with a data-type.
- Columns may have zero or more constraints.
- The data in table is in multiple rows.
- Each row have multiple values (as per columns).



Database physical layout

↳ RDBMS specific C:/ProgramData/mysql ... ↳ windows .

- In MySQL, the data is stored on disk in its data directory i.e. /var/lib/mysql
- Each database/schema is a separate sub-directory in data dir.
- Each table in the db, is a file on disk. ↳ stud.ibd
- e.g. student table in current db is stored in file /var/lib/mysql/db/student.ibd.
- Data is stored in binary format.
- A file may not be contiguously stored on hard disk.
- Data rows are not contiguous. They are scattered in the hard disk.
- In one row, all fields are consecutive.
- When records are selected, they are selected in any order.



MySQL data types

- RDBMS have similar data types (but not same).
- MySQL data types can be categorised as follows

- **Numeric types (Integers)**

- TINYINT (1 byte), SMALLINT (2 byte), MEDIUMINT (3 byte), INT (4 byte), BIGINT (8 byte), BIT(n bits)
- integer types can signed (default) or unsigned.

- **Numeric types (Floating point)**

- approx. precision – FLOAT (4 byte), DOUBLE (8 byte) | DECIMAL(m, n) – exact precision

- **Date/Time types**

duration
± 839 hrs
 $2021-10-05\ 09:52:00$ → SQL time
java.util.Date
num of seconds from
1-1-1970 00:00:00 ← epoch
1901 to
2155

- DATE, TIME, DATETIME, TIMESTAMP, YEAR

- **String types** – size = number of chars * size of char

1-1-1000
to
31-12-9999

- CHAR(1-255) – Fixed length, Very fast access. → `char or var [1@]`; 
- VARCHAR(1-65535) – Variable length, Stores length + chars. 
- TINYTEXT (255), TEXT (64K), MEDIUMTEXT (16M), LONGTEXT (4G) – Variable length, Slower access.

- **Binary types** – size = number of bytes → images, mp3, video, ...

- BINARY, VARBINARY, TINYBLOB, BLOB, MEDIUMBLOB, LONGBLOB

- **Miscellaneous types**

- ENUM, SET

→ 4 bytes = 32-bits
signed = 1 bit sign + 31 bit data = $\pm 2^{31}$
unsigned = 32 bit data = 2^{32}

CHAR vs VARCHAR vs TEXT

- CHAR

- Fixed inline storage.
- If smaller data is given, rest of space is unused.
- Very fast access.

- VARCHAR

- Variable inline storage.
- Stores length and characters.
- Slower access than CHAR.

- TEXT

- Variable external storage.
- Very slow access.
- Not ideal for indexing.


sizeof a char

- ascii = 1 byte
- unicode = 2 byte
- '
- UTF16 BE
- UTF16 LE
- UTF ...
- eabcdif = 4 byte

CHAR(4) → max 4 chars

- Ascii - 4 bytes
- Unicode - 8 bytes .

server disk



• CREATE TABLE temp(c1 CHAR(4), c2 VARCHAR(4), c3 TEXT(4));

• DESC temp;

• INSERT INTO temp VALUES('abcd', 'abcd', 'abcdef');



SQL scripts *collection of sql commands in a text file (.sql)*

- SQL script is multiple SQL queries written into a .sql file.
- SQL scripts are mainly used while database backup and restore operations.
- SQL scripts can be executed from terminal as:
 - terminal> mysql –u user –password db </path/to/sqlfile
- SQL scripts can be executed from command line as:
 - ✓ mysql> SOURCE /path/to/sqlfile
- Note that SOURCE is MySQL CLI client command.
- It reads commands one by one from the script and execute them on server.



SELECT – DQL

- Select all columns (in fixed order).
 - SELECT * FROM table;
- Select specific columns / in arbitrary order. ↗ projection
 - SELECT c1, c2, c3 FROM table;
- Column alias ↗ optional
 - SELECT c1 AS col1, c2 col2 FROM table;
- Computed columns.
 - SELECT c1, c2, c3, expr1, expr2 FROM table;
SELECT c1,
CASE WHEN condition1 THEN value1,
WHEN condition2 THEN value2,
...
ELSE valueN
END
FROM table;



SELECT – DQL

- Distinct values in column.
 - `SELECT DISTINCT c1 FROM table;`
 - `SELECT DISTINCT c1, c2 FROM table;`
- Select limited rows.
 - `SELECT * FROM table LIMIT n;`
 - `SELECT * FROM table LIMIT m, n;`

Fetch

skip *Fetch*

The diagram illustrates the execution plan for SELECT statements with LIMIT. It shows two cases: 1) LIMIT n; where the entire result set is fetched. 2) LIMIT m, n; where the first m rows are skipped and then the remaining n rows are fetched.

SELECT – DQL – WHERE

- ~~BETWEEN~~ operator (include both ends)

- c1 BETWEEN val1 AND val2

- ~~IN~~ operator (equality check with multiple values)

- c1 IN (val1, val2, val3)

- LIKE operator (similar strings)

- c1 LIKE 'pattern'.
 - % represent any number of any characters.
 - _ represent any single character.





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>





MySQL - RDBMS

Trainer: Mr. Nilesh Ghule



SELECT – DQL – WHERE

- ~~BETWEEN~~ operator (include both ends)

- c1 BETWEEN val1 AND val2

- ~~IN~~ operator (equality check with multiple values)

- c1 IN (val1, val2, val3)

- LIKE operator (similar strings)

- c1 LIKE 'pattern'.
 - % represent any number of any characters.
 - _ represent any single character.



UPDATE – DML

- To change one or more rows in a table.
- Update row(s) single column.
 - UPDATE table SET c2=new-value WHERE c1=some-value;
- Update multiple columns.
 - UPDATE table SET c2=new-value, c3=new-value WHERE c1=some-value;
- Update all rows single column.
 - UPDATE table SET c2=new-value;



DELETE – DML vs TRUNCATE – DDL vs DROP – DDL

• DELETE


- To delete one or more rows in a table.
- Delete row(s)
 - DELETE FROM table WHERE c1=value;
- Delete all rows
 - DELETE FROM table ;

• TRUNCATE

- Delete all rows.
 - TRUNCATE TABLE table;
- Truncate is faster than DELETE.

• DROP

- Delete all rows as well as table structure.
 - DROP TABLE table;
 - DROP TABLE table IF EXISTS;
- Delete database/schema.
 - DROP DATABASE db;



Seeking HELP

- HELP is client command to seek help on commands/functions.
 - HELP SELECT;
 - HELP Functions;
 - HELP SIGN;



DUAL table

- A dummy/in-memory a table having single row & single column.
- It is used for arbitrary calculations, testing functions, etc.
 - `SELECT 2 + 3 * 4 FROM DUAL;`
 - `SELECT NOW() FROM DUAL;`
 - `SELECT USER(), DATABASE() FROM DUAL;`
- In MySQL, DUAL keyword is optional.
 - `SELECT 2 + 3 * 4;`
 - `SELECT NOW();`
 - `SELECT USER(), DATABASE();`



SQL functions

- RDBMS provides many built-in functions to process the data.
- These functions can be classified as:
 - Single row functions
 - One row input produce one row output.
 - e.g. ABS(), CONCAT(), IFNULL(), ...
 - Multi-row or Group functions
 - Values from multiple rows are aggregated to single value.
 - e.g. SUM(), MIN(), MAX(), ...
- These functions can also be categorized based on data types or usage.
 - Numeric functions
 - String functions
 - Date and Time functions
 - Control flow functions
 - Information functions
 - Miscellaneous functions



Numeric & String functions

- ABS()
- POWER()
- ROUND(), FLOOR(), CEIL()

- ASCII(), CHAR()
- CONCAT()
- SUBSTRING()
- LOWER(), UPPER()
- TRIM(), LTRIM(), RTRIM()
- LPAD(), RPAD()
- REGEXP_LIKE()





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>



MySQL - RDBMS

Trainer: Mr. Nilesh Ghule



SQL functions

- RDBMS provides many built-in functions to process the data.
- These functions can be classified as:
 - Single row functions
 - One row input produce one row output.
 - e.g. ABS(), CONCAT(), IFNULL(), ...
 - Multi-row or Group functions
 - Values from multiple rows are aggregated to single value.
 - e.g. SUM(), MIN(), MAX(), ...
- These functions can also be categorized based on data types or usage.
 - Numeric functions
 - String functions
 - Date and Time functions
 - Control flow functions
 - Information functions
 - Miscellaneous functions



Numeric & String functions

- ABS()
- POWER()
- ROUND(), FLOOR(), CEIL()

- ASCII(), CHAR()
- CONCAT()
- SUBSTRING()
- LOWER(), UPPER()
- TRIM(), LTRIM(), RTRIM()
- LPAD(), RPAD()
- REGEXP_LIKE()



Date-Time and Information functions

- VERSION()
- USER(), DATABASE()
- MySQL supports multiple date time related data types
 - DATE (3), TIME (3), DATETIME (5), TIMESTAMP (4), YEAR (1)
- SYSDATE(), NOW()
- DATE(), TIME()
- DAYOFMONTH(), MONTH(), YEAR(), HOUR(), MINUTE(), SECOND(), ...
- DATEDIFF(), DATE_ADD(), TIMEDIFF()
- MAKEDATE(), MAKETIME()




Control and NULL and List functions

- NULL is special value in RDBMS that represents absence of value in that column.
 - NULL values do not work with relational operators and need to use special operators.
 - Most of functions return NULL if NULL value is passed as one of its argument.
 - ISNULL()
 - IFNULL()
 - NULLIF()
 - COALESCE()
-
- GREATEST(), LEAST()
 - IF(condition, true-value, false-value)




Group functions



- Work on group of rows of table.
- Input to function is data from multiple rows & then output is single row. Hence these functions are called as "Multi Row Function" or "Group Functions".
- These functions are used to perform aggregate ops like sum, avg, max, min, count or std dev, etc. Hence these fns are also called as "Aggregate Functions".
- Example: SUM(), AVG(), MAX(), MIN(), COUNT().
- NULL values are ignored by group functions.
- Limitations of GROUP functions:
 - Cannot select group function along with a column.
 - Cannot select group function along with a single row fn.
 - Cannot use group function in WHERE clause/condition.
 - Cannot nest a group function in another group fn.




GROUP BY clause

- GROUP BY is used for analysis of data i.e. generating reports & charts.
- When GROUP BY single column, generated output can be used to plot 2-D chart.
When GROUP BY two column, generated output can be used to plot 3-D chart and so on.
- GROUP BY queries are also called as Multi-dimensional / Spatial queries.
- Syntactical Characteristics:
 - If a column is used for GROUP BY, then it may or may not be used in SELECT clause.
 - If a column is in SELECT, it must be in GROUP BY.
- When GROUP BY query is fired on database server, it does following:
 - ✓ Load data from server disk into server RAM.
 - ✓ Sort data on group by columns.
 - ✓ Group similar records by group columns.
 - ✓ Perform given aggregate ops on each column.
 - ✓ Send result to client.







HAVING clause

- HAVING clause cannot be used without GROUP BY clause.
- HAVING clause is used to specify condition on aggregate values.
- Examples:
 - SELECT deptno, SUM(sal) FROM EMP GROUP BY deptno HAVING SUM(sal) > 9000;
- Syntactical Characteristics:
 - WHERE clause executed for each record; while HAVING is executed for each group.
 - HAVING clause can be used to specify condition on group fn or grouped columns.
 - However using HAVING to specify condition of group col reduce the performance. Use WHERE clause for the same.
- Examples:
 - SELECT deptno, SUM(sal) FROM EMP GROUP BY deptno HAVING deptno = 20; ✓
 - SELECT deptno, SUM(sal) FROM EMP WHERE deptno = 20 GROUP BY deptno; ✓
- We may use GROUP BY with WHERE, ORDER BY & LIMIT. ✓





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>






MySQL - RDBMS

Trainer: Mr. Nilesh Ghule




Entity Relations

- To avoid redundancy of the data, data should be organized into multiple tables so that tables are related to each other.
- The relations can be one of the following
 - One to One
 - One to Many
 - Many to One
 - Many to Many
- Entity relations is outcome of Normalization process.



SQL Joins

- Join statements are used to SELECT data from multiple tables using single query.
- Typical RDBMS supports following types of joins:
 - Cross Join
 - Inner Join
 - Left Outer Join
 - Right Outer Join
 - Full Outer Join
 - Self join



Cross Join

dept 4

deptno	dname
10	DEV
20	QA
30	OPS
40	ACC

emp 5

empno	ename
1	Amit
2	Rahul
3	Nilesh
4	Nitin
5	Sarang

```

for (Emp e: emp) { driving table (outer loop)
    for (Dept d: dept) { driven table (inner loop)
        point(e.eename, d.dname);
    }
}

```

3

✓ Select e.ename, d.dname from
emp e cross join depts d;

- Compares each row of Table1 with every row of Table2.
 - Yields all possible combinations of Table1 and Table2.
 - In MySQL, The larger table is referred as "Driving Table", while smaller table is referred as "Driven Table". Each row of Driving table is combined with every row of Driven table.
 - Cross join is the fastest join, because there is no condition check involved.



Inner Join

dept (4)

deptno	dname
10	DEV
20	QA
30	OPS
40	ACC


emp (5)

empno	ename	deptno
1	Amit	10
2	Rahul	10
3	Nilesh	20
4	Nitin	50
5	Sarang	50



Amit Dev
Rahul Dev
Nilesh QA

QA
OPS
Acc




Select e.ename, d.dname from empss e
inner join dept d on e.deptno = d.deptno;

```
for(Emp e: emp) {  
    for(Dept d : dept) {  
        if(e.deptno == d.deptno)  
            point(e.ename, d.dname);  
    }  
}
```

- The inner JOIN is used to return rows from both tables that satisfy the join condition.
- Non-matching rows from both tables are skipped.
- If join condition contains equality check, it is referred as equi-join; otherwise it is non-equi-join.

Left Outer Join

deptno	dname
10	DEV
20	QA
30	OPS
40	ACC



empno	ename	deptno
1	Amit	10
2	Rahul	10
3	Nilesh	20
4	Nitin	50
5	Sarang	50

Select e.ename, d.dname
From dept d
left outer join emp e
on d.deptno = e.deptno;


```

for(Dept d : dept) {
    found=false;
    for(Emp e : emp) {
        if(e.deptno == d.deptno) {
            point(e.ename,d.dname)
            found=true;
        }
    }
    if(found == false)
        point(NULL,d.dname);
}
  
```

- Left outer join is used to return matching rows from both tables along with additional rows in left table.
- Corresponding to additional rows in left table, right table values are taken as NULL.
- OUTER keyword is optional.

Right Outer Join

deptno	dname
10	DEV
20	QA
30	OPS
40	ACC



empno	ename	deptno
1	Amit	10
2	Rahul	10
3	Nilesh	20
4	Nitin	50
5	Sarang	50

select e.ename, d.dname
from depts d
right outer join emps e
on d.deptno = e.deptno;

- Right outer join is used to return matching rows from both tables along with additional rows in right table.
- Corresponding to additional rows in right table, left table values are taken as NULL.
- OUTER keyword is optional.

Full Outer Join

deptno	dname
10	DEV
20	QA
30	OPS
40	ACC

empno	ename	deptno
1	Amit	10
2	Rahul	10
3	Nilesh	20
4	Nitin	50
5	Sarang	50

- Full join is used to return matching rows from both tables along with additional rows in both tables.
- Corresponding to additional rows in left or right table, opposite table values are taken as NULL.
- Full outer join is not supported in MySQL, but can be simulated using set operators.

Set operators

query1 output + query2 output = *union or union all*

ename	dname
Amit	DEV
Rahul	DEV
Nilesh	QA
NULL	OPS
NULL	ACC

join1

ename	dname
Amit	DEV
Rahul	DEV
Nilesh	QA
Nitin	NULL
Sarang	NULL

join2

num of columns must be same.

- UNION operator is used to combine results of two queries. The common data is taken only once. It can be used to simulate full outer join.
- UNION ALL operator is used to combine results of two queries. Common data is repeated.



Self Join

- When join is done on same table, then it is known as "Self Join". The both columns in condition belong to the same table.
- Self join may be an inner join or outer join.

empno	ename	deptno	mgr
1	Amit	10	4
2	Rahul	10	3
3	Nilesh	20	4
4	Nitin	50	5
5	Sarang	50	NULL


emps e

empno	ename	deptno
1	Amit	10
2	Rahul	10
3	<u>Nilesh</u>	20
4	<u>Nitin</u>	50
5	Sarang	50

emps m

select e.ename, m.ename from emps e
inner join emps m on e.mgr = m.empno;
select e.ename, m.ename from emps e
left join emps m on e.mgr = m.empno;

Multi-Table Joins





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>



MySQL - RDBMS

Trainer: Mr. Nilesh Ghule



Sub queries

- Sub-query is query within query. Typically it work with SELECT statements.
- Output of inner query is used as input to outer query.
- If no optimization is enabled, for each row of outer query result, sub-query is executed once. This reduce performance of sub-query.
- Single row sub-query
 - Sub-query returns single row.
 - Usually it is compared in outer query using relational operators.



Sub queries

- Multi-row sub-query
 - Sub-query returns multiple rows.
 - Usually it is compared in outer query using operators like IN, ANY or ALL.
 - IN operator compare for equality with results from sub-queries (at least one result should match)..
 - ANY operator compares with the results from sub-queries (at least one result should match).
 - ALL operator compares with the results from sub-queries (all results should match).
- GR } { AND



Sub queries

- Correlated sub-query

- If number of results from sub-query are reduced, query performance will increase.
- This can be done by adding criteria (WHERE clause) in sub-query based on outer query row.
- Typically correlated sub-query use IN, ALL, ANY and EXISTS operators.

EXISTS → inner query returns 1 or more rows.

NOT EXISTS → inner query returns 0 rows.



Sub query

- Sub queries with UPDATE and DELETE are not supported in all RDBMS.
 - In MySQL, Sub-queries in UPDATE/DELETE is allowed, but sub-query should not SELECT from the same table, on which UPDATE/DELETE operation is in progress.
-



Views

- RDBMS view represents view (projection) of the data.
- View is based on SELECT statement.
- Typically it is restricted view of the data (limited rows or columns) from one or more tables (joins and/or sub-queries) or summary of the data (grouping).
- Data of view is not stored on server hard-disk; but its SELECT statement is stored in compiled form. It speed up execution of view.



Views

- Views are of two types: Simple view and Complex view
- Usually if view contains computed columns, group by, joins or sub-queries, then the views are said to be complex. DML operations are not supported on these views.
- DML operations on view affects underlying table.
- View can be created with CHECK OPTION to ensure that DML operations can be performed only the data visible in that view.



View

- Views can be differentiated with: SHOW FULL TABLES.
 - Views can be dropped with DROP VIEW statement.
 - View can be based on another view.
-
- Applications of views
 - Security: Providing limited access to the data.
 - Hide source code of the table.
 - Simplifies complex queries.





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>






MySQL - RDBMS

Trainer: Mr. Nilesh Ghule



Transaction

- Transaction is set of DML queries executed as a single unit.
- Transaction examples
 - accounts table [id, type, balance]
 - UPDATE accounts SET balance=balance-1000 WHERE id = 1;
 - UPDATE accounts SET balance=balance+1000 WHERE id = 2;
- RDBMS transaction have ACID properties.
 - Atomicity
 - All queries are executed as a single unit. If any query is failed, other queries are discarded.
 - Consistency
 - When transaction is completed, all clients see the same data.
 - Isolation
 - Multiple transactions (by same or multiple clients) are processed concurrently.
 - Durable
 - When transaction is completed, all data is saved on disk.



Transaction

- Transaction management
 - START TRANSACTION;
 - ...
 - COMMIT WORK;
- START TRANSACTION;
 - ...
 - ROLLBACK WORK;
- In MySQL autocommit variable is by default 1. So each DML command is auto-committed into database.
 - SELECT @@autocommit;
- Changing autocommit to 0, will create new transaction immediately after current transaction is completed. This setting can be made permanent in config file.
 - SET autocommit=0;



Views

- RDBMS view represents view (projection) of the data.
- View is based on SELECT statement.
- Typically it is restricted view of the data (limited rows or columns) from one or more tables (joins and/or sub-queries) or summary of the data (grouping).
- Data of view is not stored on server hard-disk; but its SELECT statement is stored in compiled form. It speed up execution of view.



Views

- Views are of two types: Simple view and Complex view
- Usually if view contains computed columns, group by, joins or sub-queries, then the views are said to be complex. DML operations are not supported on these views.
- DML operations on view affects underlying table.
- View can be created with CHECK OPTION to ensure that DML operations can be performed only the data visible in that view.




View

- Views can be differentiated with: SHOW FULL TABLES.
- Views can be dropped with DROP VIEW statement.
- View can be based on another view.

- Applications of views

- Security: Providing limited access to the data.
- Hide source code of the table.
- Simplifies complex queries.

Example joins/subqueries are added into view - select.



Data Control Language

- Security is built-in feature of any RDBMS. It is implemented in terms of permissions (a.k.a. privileges).
- There are two types of privileges.
- System privileges
 - Privileges for certain commands i.e. CREATE, ALTER, DROP, ...
 - Typically these privileges are given to the database administrator or higher authority user.
- Object privileges
 - RDBMS objects are table, view, stored procedure, function, triggers, ...
 - Can perform operations on the objects i.e. INSERT, UPDATE, DELETE, SELECT, CALL, ...
 - Typically these privileges are given to the database users.





User Management

- User management is responsibility of admin (root).
- New user can be created using CREATE USER.
 - ✓ CREATE USER user@host IDENTIFIED BY 'password';
 - host can be hostname of server, localhost (current system) or '%' for all client systems.
- Permissions for the user can be listed using SHOW GRANTS command.
 - ✓ SHOW GRANTS FOR user@host;
- Users can be deleted using DROP USER.
 - ✓ DROP USER 'user'@'host';
- Change user password.
 - ALTER USER user@host IDENTIFIED BY 'new_password';
 - FLUSH PRIVILEGES; *To reload security settings if they are modified using DML ops on System tables.*

In RDBMS, all config/settings/meta data is stored in system tables/databases.

mysql.user
db table



Data Control Language

everything in db.

- Permissions are given to user using GRANT command.
 - GRANT CREATE ON db.* TO user@host;
 - GRANT CREATE ON *.* TO user1@host, ~~user2@host~~;
 - GRANT SELECT ON db.table TO user@host;
 - GRANT SELECT, INSERT, UPDATE ON db.table TO user@host;
 - GRANT ALL ON db.* TO user@host;
- By default one user cannot give permissions to other user. This can be enabled using **WITH GRANT OPTION**.
 - GRANT ALL ON *.* TO user@host WITH GRANT OPTION;
- Permissions assigned to any user can be withdrawn using REVOKE command.
 - REVOKE SELECT, INSERT ON db.table FROM user@host;
- Permissions can be activated by FLUSH PRIVILEGES.
 - System GRANT tables are reloaded by this command. Auto done after GRANT, REVOKE.
 - Command is necessary if GRANT tables are modified using DML operations.





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>




MySQL - RDBMS

Trainer: Mr. Nilesh Ghule




Transaction

- Transaction is set of DML queries executed as a single unit.
- Transaction examples
 - accounts table [id, type, balance]
 - UPDATE accounts SET balance=balance-1000 WHERE id = 1;
 - UPDATE accounts SET balance=balance+1000 WHERE id = 2;
- RDBMS transaction have ACID properties.
 - Atomicity
 - All queries are executed as a single unit. If any query is failed, other queries are discarded.
 - Consistency
 - When transaction is completed, all clients see the same data.
 - Isolation
 - Multiple transactions (by same or multiple clients) are processed concurrently.
 - Durable
 - When transaction is completed, all data is saved on disk.



Transaction

- Transaction management
 - START TRANSACTION;
 - ...
 - COMMIT WORK;
 - START TRANSACTION;
 - ...
 - ROLLBACK WORK;
- In MySQL autocommit variable is by default 1. So each DML command is auto-committed into database.
 - SELECT @@autocommit; ↗ 
- Changing autocommit to 0, will create new transaction immediately after current transaction is completed. This setting can be made permanent in config file.
 - SET autocommit=0;



Transaction

- Save-point is state of database tables (data) at the moment (within a transaction).
- It is advised to create save-points at end of each logical section of work.
- Database user may choose to rollback to any of the save-point.
- Transaction management with Save-points
 - START TRANSACTION;
 - ...
 - SAVEPOINT sa1;
 - ...
 - SAVEPOINT sa2;
 - ...
 - ROLLBACK TO sa1;
 - ...
 - COMMIT; // or ROLLBACK
- Commit always commit the whole transaction.
- ROLLBACK or COMMIT clears all save-points.




Transaction

- Transaction is set of DML statements.
- If any DDL statement is executed, current transaction is automatically committed.
- Any power failure, system or network failure automatically rollback current state.
- Transactions are isolated from each other and are consistent.



Row locking

- When an user update or delete a row (within a transaction), that row is locked and becomes read-only for other users.
- The other users see old row values, until transaction is committed by first user.
- If other users try to modify or delete such locked row, their transaction processing is blocked until row is unlocked.
- Other users can INSERT into that table. Also they can UPDATE or DELETE other rows.
- The locks are automatically released when COMMIT/ROLLBACK is done by the user. *or timeout occurs.*
- This whole process is done automatically in MySQL. It is called as "OPTIMISTIC LOCKING".




Row locking

- Manually locking the row in advanced before issuing UPDATE or DELETE is known as "PESSIMISTIC LOCKING".
- This is done by appending FOR UPDATE to the SELECT query.
- It will lock all selected rows, until transaction is committed or rolled back.
or timeout.
- If these rows are already locked by another users, the SELECT operation is blocked until rows lock is released.
- By default MySQL does table locking.
Row locking is possible only when table is indexed on the column. *(pri key)*



Index

- Index enable faster searching in tables by indexed columns.
 - CREATE INDEX idx_name ON table(column);
- One table can have multiple indexes on different columns/order.
- Typically indexes are stored as some data structure (like BTREE or HASH) on disk.
- Indexes are updated during DML operations. So DML operation are slower on indexed tables.



Index

- Index can be ASC or DESC.
 - It cause storage of key values in respective order (MySQL 8.x onwards).
 - ASC/DESC index is used by optimizer on ORDER BY queries.
- There are four types of indexes:
 - Simple index
 - CREATE INDEX idx_name ON table(column [ASC|DESC]);
 - Unique index
 - CREATE UNIQUE INDEX idx_name ON table(column [ASC|DESC]);
 - Doesn't allow duplicate values.
 - Composite index
 - CREATE INDEX idx_name ON table(column1 [ASC|DESC], column2 [ASC|DESC]);
 - Composite index can also be unique. Do not allow duplicate combination of columns.
 - Clustered index
 - PRIMARY index automatically created on Primary key for row lookup.
 - If primary key is not available, hidden index is created on synthetic column.
 - It is maintained in tabular form and its reference is used in other indexes.



Index

- Indexes should be created on shorter (INT, CHAR, ...) columns to save disk space.
- Few RDBMS do not allow indexes on external columns i.e. TEXT, BLOB.
- MySQL support indexing on TEXT/BLOB up to n characters.
 - CREATE TABLE test (blob_col BLOB, ..., INDEX(blob_col(10)));
- To list all indexes on table:
 - SHOW INDEXES FROM table;
- To drop an index:
 - DROP INDEX idx_name ON table;
- When table is dropped, all indexes are automatically dropped.
- Indexes should not be created on the columns not used frequent search, ordering or grouping operations.
- Columns in join operation should be indexed for better performance.



Query performance

- Few RDBMS features ensure better query performance.
 - Index speed up execution of SELECT queries (search operations).
 - Correlated sub-queries execute faster.
- Query performance can be observed using EXPLAIN statement.
 - EXPLAIN FORMAT=JSON SELECT ...;
- EXPLAIN statement shows
 - Query cost (Lower is the cost, faster is the query execution).
 - Execution plan (Algorithm used to execute query e.g. loop, semi-join, materialization, etc).
- Optimizations can be enabled or disabled by optimizer_switch system variable.
 - SELECT @@optimizer_switch;
 - SET @@optimizer_switch='materialization=off';



Constraints

- Constraints are restrictions imposed on columns.
- There are five constraints
 - NOT NULL
 - UNIQUE
 - PRIMARY KEY
 - FOREIGN KEY
 - CHECK
- Few constraints can be applied at either column level or table level. Few constraints can be applied on both.
- Optionally constraint names can be mentioned while creating the constraint. If not given, it is auto-generated.
- Each DML operation check the constraints before manipulating the values. If any constraint is violated, error is raised.





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>





MySQL - RDBMS

Trainer: Mr. Nilesh Ghule



Constraints

- Constraints are restrictions imposed on columns.
- There are five constraints
 - NOT NULL
 - UNIQUE
 - PRIMARY KEY
 - FOREIGN KEY
 - CHECK
- Few constraints can be applied at either column level or table level. Few constraints can be applied on both.
- Optionally constraint names can be mentioned while creating the constraint. If not given, it is auto-generated.
- Each DML operation check the constraints before manipulating the values. If any constraint is violated, error is raised.



Constraints

- **PRIMARY KEY**

- Column or set of columns that uniquely identifies a row.
- Only one primary key is allowed for a table.
- Primary key column cannot have duplicate or NULL values.
- Internally index is created on PK column.
- TEXT/BLOB cannot be primary key.
- If no obvious choice available for PK, composite or surrogate PK can be created.
- Creating PK for a table is a good practice.
- PK can be created at table level or column level.
- CREATE TABLE table(c1 TYPE PRIMARY KEY, ...);
- CREATE TABLE table(c1 TYPE, ..., PRIMARY KEY(c1));
- CREATE TABLE table(c1 TYPE, ..., CONSTRAINT constraint_name PRIMARY KEY(c1));
- CREATE TABLE table(c1 TYPE, c2 TYPE, ..., PRIMARY KEY(c1, c2));



Constraints

- FOREIGN KEY

- Column or set of columns that references a column of some table.
- If column belongs to the same table, it is "self referencing".
- Foreign key constraint is specified on child table column.
- FK can have duplicate values as well as null values.
- FK constraint is applied on column of child table (not on parent table).
- Child rows cannot be deleted, until parent rows are deleted.
- MySQL have ON DELETE CASCADE clause to ensure that child rows are automatically deleted, when parent row is deleted. ON UPDATE CASCADE clause does same for UPDATE operation.
- By default foreign key checks are enabled. They can be disabled by
 - `SET @@foreign_key_checks = 0;`
- FK constraint can be applied on table level as well as column level.
- `CREATE TABLE child(c1 TYPE, ..., FOREIGN KEY (c1) REFERENCES parent(col))`



Constraints

- **CHECK**
 - CHECK is integrity constraint in SQL.
 - CHECK constraint specifies condition on column.
 - Data can be inserted/updated only if condition is true; otherwise error is raised.
 - CHECK constraint can be applied at table level or column level.
 - `CREATE TABLE table(c1 TYPE, c2 TYPE CHECK condition1, ..., CHECK condition2);`



DDL – ALTER statement

- ALTER statement is used to do modification into table, view, function, procedure, ...
- ALTER TABLE is used to change table structure.
- Add new column(s) into the table.
 - ALTER TABLE table ADD col TYPE;
 - ALTER TABLE table ADD c1 TYPE, c2 TYPE;
- Modify column of the table.
 - ALTER TABLE table MODIFY col NEW_TYPE;
- Rename column.
 - ALTER TABLE CHANGE old_col new_col TYPE;
- Drop a column
 - ALTER TABLE ^{table name} DROP COLUMN col;
- Rename table
 - ALTER TABLE table RENAME TO new_table;




MySQL Programming

- RDBMS Programming is an ISO standard – part of SQL standard – since 1992.
- SQL/PSM stands for Persistent Stored Module.
- Inspired from PL/SQL - Programming language of Oracle.
- PSM allows writing programs for RDBMS. The program contains set of SQL statements along with programming constructs e.g. variables, if-else, loops, case, ...
- PSM is a block language. Blocks can be nested into another block.
- MySQL program can be a stored procedure, function or trigger.



MySQL Programming

- MySQL PSM program is written by db user (programmers).
- It is submitted from client, server check syntax & store them into db in compiled form.
- The program can be executed by db user when needed.
- Since programs are stored on server in compiled form, their execution is very fast.
- All these programs will run in server memory.



Stored Procedure

- Stored Procedure is a routine. It contains multiple SQL statements along with programming constructs. *→ no RETURN value Keyword.*
- Procedure doesn't return any value (like void fns in C).
- Procedures can take zero or more parameters.
- Procedures are created using CREATE PROCEDURE and deleted using DROP PROCEDURE.
- Procedures are invoked/called using CALL statement.
- Result of stored procedure can be
 - returned via OUT parameter.
 - inserted into another table.
 - produced using SELECT statement (at end of SP).
- Delimiter should be set before writing procedure.



Stored Procedure

```
CREATE TABLE result(v1 DOUBLE, v2 VARCHAR(50));          -- 01_hello.sql (using editor)
DELIMITER $$                                           DROP PROCEDURE IF EXISTS sp_hello;
                                                       DELIMITER $$

CREATE PROCEDURE sp_hello()                           CREATE PROCEDURE sp_hello()
BEGIN                                              BEGIN
    INSERT INTO result VALUES(1, 'Hello World');   SELECT 1 AS v1, 'Hello World' AS v2;
END;                                                 END;

$$                                                 $$

DELIMITER ;                                         DELIMITER ;
                                                       SOURCE /path/to/01_hello.sql
CALL sp_hello();                                     CALL sp_hello();

SELECT * FROM result;
```



Stored Procedure – PSM Syntax

VARIABLES

```
DECLARE varname DATATYPE;  
DECLARE varname DATATYPE DEFAULT init_value;  
SET varname = new_value;  
SELECT new_value INTO varname;  
SELECT expr_or_col INTO varname FROM table_name;
```

PARAMETERS

```
CREATE PROCEDURE sp_name(PARAMTYPE p1 DATATYPE)  
BEGIN  
...  
END;  
  
-- IN param: Initialized by calling program.  
-- OUT param: Initialized by called procedure.  
-- INOUT param: Initialized by calling program and  
modified by called procedure  
-- OUT & INOUT param declared as session variables.  
  
CREATE PROCEDURE sp_name(OUT p1 INT)  
BEGIN  
    SELECT 1 INTO p1;  
END;  
  
SET @res = 0;  
CALL sp_name(@res);  
SELECT @res;
```

IF-ELSE

```
IF condition THEN  
    body;  
END IF;  
-----  
IF condition THEN  
    if-body;  
ELSE  
    else-body;  
END IF;  
-----  
IF condition THEN  
    if1-body;  
ELSE  
    IF condition THEN  
        if2-body;  
    ELSE  
        else2-body;  
    END IF;  
END IF;  
-----  
IF condition THEN  
    if1-body;  
ELSEIF condition THEN  
    if2-body;  
ELSE  
    else-body;  
END IF;
```

LOOPS

```
WHILE condition DO  
    body;  
END WHILE;  
-----  
REPEAT  
    body;  
UNTIL condition  
END REPEAT;  
-----  
label: LOOP  
IF condition THEN  
    ...  
    LEAVE label;  
END IF;  
...  
END LOOP;
```

SHOW PROCEDURE

```
SHOW PROCEDURE STATUS  
LIKE 'sp_name';  
  
SHOW CREATE PROCEDURE sp_name;
```

DROP PROCEDURE

```
DROP PROCEDURE  
IF EXISTS sp_name;
```

CASE-WHEN

```
CASE  
WHEN condition THEN  
    body;  
WHEN condition THEN  
    body;  
ELSE  
    body;  
END CASE;
```





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>





MySQL - RDBMS

Trainer: Mr. Nilesh Ghule



Normalization

- Concept of table design: Table, Structure, Data Types, Width, Constraints, Relations.
- Goals:
 - Efficient table structure.
 - Avoid data redundancy i.e. unnecessary duplication of data (to save disk space).
 - Reduce problems of insert, update & delete.
*insert anomaly
update anomaly
delete anomaly*
- Done from input perspective.
- Based on user requirements.
- Part of software design phase.
- View entire appln on per transaction basis & then normalize each transaction separately.
- Transaction Examples:
 - Banking, Rail Reservation, Online Shopping.



Normalization

- For given transaction make list of all the fields.
- Strive for atomicity.
- Get general description of all field properties.
- For all practical purposes we can have a single table with all the columns. Give meaningful names to the table.
- Assign datatypes and widths to all columns on the basis of general desc of fields properties.
- Remove computed columns.
- Assign primary key to the table. ?
- At this stage data is in un-normalized form.
- UNF is starting point of normalization.



SQL Keys

- An SQL key is either a single column (or attribute) or a group of columns that can uniquely identify rows (or tuples) in a table.
- Super key is a single key or a group of multiple keys that can uniquely identify tuples in a table.
- Candidate key is a single key or a group of multiple keys that uniquely identify rows in a table.
- Primary key is the Candidate key selected by the database administrator to uniquely identify tuples in a table.
- Alternate keys are those candidate keys which are not the Primary key.
- Foreign key is an attribute which is a Primary key in its parent table, but is included as an attribute in another related table.



SQL Keys

empno	ename	address	dept	job	ssn/aadhar	passport	email	phone
1	A	PUNE	10	ANALYST	111100001111	11111111	a@sun.com	9822012345
2	B	DELHI	10	CLERK	222200002222	22222222	b@sun.com	9422012345
3	C	MUMBAI	30	PRESIDENT	333300003333	33333333	c@sun.com	9822123456
4	D	CHENNAI	20	MANAGER	444400004444	44444444	d@sun.com	9422123456
5	E	GOA	20	CLERK	555500005555	55555555	e@sun.com	9822123321

S	S	S	S	S	S	S	S
C	C	C	C	C	C	C	C
PK			AK	AK	AK	AK	AK



Normalization

- 1. Remove repeating group into a new table.
- 2. Key elements will be PK of new table.
- 3. (Optional) Add PK of original table to new table to give us Composite PK.
 - Repeat steps 1-3 infinitely -- to remove all repeating groups into new tables.
 - This is **1-NF**. No repeating groups present here. One to Many relationship between two tables.



Normalization

- 4. Only table with composite PK to be examined.
- 5. Those columns that are not dependent on the entire composite PK, they are to be removed into a new table.
- 6. The key elements on which the non-key elements were originally dependent, it is to be added to the new table, and it will be the PK of new table.
 - Repeat steps 4-6 infinitely -- to separate all non-key elements from all tables with composite primary key.
 - This is **2-NF**. Many-to-Many relationship.



Normalization

- 7. Only non-key elements are examined for inter-dependencies.
- 8. Inter-dependent cols that are not directly related to PK, they are to be removed into a new table.
- 9. (a) Key ele will be PK of new table.
- 9. (b) The PK of new table is to be retained in original table for relationship purposes.
 - Repeat steps 7-9 infinitely to examine all non-key eles from all tables and separate them into new table if not dependent on PK.
 - This is **3-NF**.



Normalization

- To ensure data consistency (no wrong data entered by end user).
- Separate table to be created of well-known data. So that min data will be entered by the end user.
- This is BCNF or 4-NF.



De-normalization

- Normalization will yield a structure that is non-redundant.
- Having too many inter-related tables will lead to complex and inefficient queries.
- To ensure better performance of analytical queries, few rules of normalization can be compromised.
- This process is de-normalization.



Codd's rules

- Proposed by Edgar F. Codd – pioneer of the RDBMS – in 1980.
- If any DBMS follow these rules, it can be considered as RDBMS.
- The 0th rule is the main rule known as “The foundation rule”.
 - For any system that is advertised as, or claimed to be, a relational data base management system, that system must be able to manage data bases entirely through its relational capabilities.
- The rest of rules can be considered as elaboration of this foundation rule.



Codd's rules

- Rule 1: The information rule:
 - All information in a relational data base is represented explicitly at the logical level and in exactly one way – by values in tables.
- Rule 2: The guaranteed access rule:
 - Each and every datum (atomic value) in a relational data base is guaranteed to be logically accessible by resorting to a combination of table name, primary key value and column name.
- Rule 3: Systematic treatment of null values:
 - Null values (distinct from the empty character string or a string of blank characters and distinct from zero or any other number) are supported in fully relational DBMS for representing missing information and inapplicable information in a systematic way, independent of data type.



Codd's rules

- Rule 4: Dynamic online catalog based on the relational model:
 - The data base description is represented at the logical level in the same way as ordinary data, so that authorized users can apply the same relational language to its interrogation as they apply to the regular data.
- Rule 5: The comprehensive data sublanguage rule:
 - A relational system may support several languages. However, there must be at least one language that supports all functionalities of a RDBMS i.e. data definition, data manipulation, integrity constraints, transaction management, authorization.



Codd's rules

- Rule 6: The view updating rule:
 - All views that are theoretically updatable are also updatable by the system.
- Rule 7: Possible for high-level insert, update, and delete:
 - The capability of handling a base relation or a derived relation as a single operand applies not only to the retrieval of data but also to the insertion, update and deletion of data.
- Rule 8: Physical data independence:
 - Application programs and terminal activities remain logically unbroken whenever any changes are made in either storage representations or access methods.
- Rule 9: Logical data independence:
 - Application programs & terminal activities remain logically unbroken when information-preserving changes of any kind that theoretically permit un-impairment are made to the base tables.



Codd's rules

- Rule 10: Integrity independence:
 - Integrity constraints specific to a particular relational database must be definable in the relational data sublanguage and storable in the catalog, not in the application programs.
- Rule 11: Distribution independence:
 - The end-user must not be able to see that the data is distributed over various locations. Users should always get the impression that the data is located at one site only.
- Rule 12: The non-subversion rule:
 - If a relational system has a low-level (single-record-at-a-time) language, that low level cannot be used to subvert or bypass the integrity rules and constraints expressed in the higher level relational language (multiple-records-at-a-time).



MySQL Exceptions

not compile time errors

- Exceptions are runtime problems, which may arise during execution of stored procedure, function or trigger.
- Required actions should be taken against these errors.
- SP execution may be continued or stopped after handling exception.
- MySQL error handlers are declared as:
 - `DECLARE action HANDLER FOR condition handler_impl;`
- The *action* can be: CONTINUE or EXIT.
- The *condition* can be:
 - MySQL error code: e.g. 1062 for duplicate entry.
 - SQLSTATE value: e.g. 23000 for duplicate entry, NOTFOUND for end-of-cursor.
 - Named condition: e.g. `DECLARE duplicate_entry CONDITION FOR 1062;`
- The *handler_impl* can be: Single liner or PSM block i.e. BEGIN ... END;



MySQL Stored Functions

- Stored Functions are MySQL programs like stored procedures.
- Functions can be having zero or more parameters. MySQL allows only IN params.
- Functions must return some value using RETURN statement.
- Function entire code is stored in system table.
- Like procedures, functions allows statements like local variable declarations, if-else, case, loops, etc. One function can invoke another function/procedure and vice-versa. The functions can also be recursive.
- There are two types of functions: DETERMINISTIC and NOT DETERMINISTIC.

CREATE FUNCTION

```
CREATE FUNCTION fn_name(p1 TYPE)
RETURNS TYPE
[NOT] DETERMINISTIC
BEGIN
    body;
    RETURN value;
END;
```

SHOW FUNCTION

```
SHOW FUNCTION STATUS LIKE 'fn_name';
SHOW CREATE FUNCTION fn_name;
```

SELECT fn_name(...);

DROP FUNCTION

```
DROP FUNCTION IF EXISTS fn_name;
```



MySQL Triggers

- Triggers are supported by all standard RDBMS like Oracle, MySQL, etc.
- Triggers are not supported by WEAK RDBMS like MS-- Access.
- Triggers are not called by client's directly, so they don't have args & return value.
- Trigger execution is caused by DML operations on database.
 - BEFORE/AFTER INSERT, BEFORE/AFTER UPDATE, BEFORE/AFTER DELETE.
- Like SP/FN, Triggers may contain SQL statements with programming constructs. They may also call other SP or FN.
- However COMMIT/ROLLBACK is not allowed in triggers. They are executed in same transaction in which DML query is executed.

CREATE TRIGGER

```
CREATE TRIGGER trig_name
AFTER|BEFORE dml_op ON table
FOR EACH ROW
BEGIN
    body;
    -- use OLD & NEW keywords
    -- to access old/new rows.
    -- INSERT triggers - NEW rows.
    -- DELETE triggers - OLD rows.
END;
```

SHOW TRIGGERS

```
SHOW TRIGGERS FROM db_name;
```

DROP TRIGGER

```
DROP TRIGGER trig_name;
```



MySQL Triggers

- Applications of triggers:
 - Maintain logs of DML operations (Audit Trails).
 - Data cleansing before insert or update data into table. (Modify NEW value).
 - Copying each record AFTER INSERT into another table to maintain "Shadow table".
 - Copying each record AFTER DELETE into another table to maintain "History table".
 - Auto operations of related tables using cascading triggers.
- Cascading triggers
 - One trigger causes execution of 2nd trigger, 2nd trigger causes execution of 3rd trigger and so on.
 - In MySQL, there is no upper limit on number of levels of cascading.
 - This is helpful in complicated business processes.
- Mutating table error
 - If cascading trigger causes one of the earlier trigger to re-execute, "mutating table" error is raised.
 - This prevents infinite loop and also rollback the current transaction.





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>



MySQL - RDBMS

Trainer: Mr. Nilesh Ghule



Sunbeam Infotech

www.sunbeaminfo.com



Sunbeam Infotech

www.sunbeaminfo.com

DBMS

- Any enterprise application need to manage data.
- In early days of software development, programmers store data into files and does operation on it. However data is highly application specific.
- Even today many software manage their data in custom formats e.g. Tally, Address book, etc.
- As data management became more common, DBMS systems were developed to handle the data. This enabled developers to focus on the business logic e.g. FoxPro, DBase, Excel, etc.
- At least CRUD (Create, Retrieve, Update and Delete) operations are supported by all databases.
- Traditional databases are file based, less secure, single-user, non-distributed, manage less amount of data (MB), complicated relation management, file-locking and need number of lines of code to use in applications.




Sunbeam Infotech

www.sunbeaminfo.com

RDBMS

- RDBMS is relational DBMS.
- It organizes data into Tables, rows and columns. The tables are related to each other.
- RDBMS follow table structure, more secure, multi-user, server-client architecture, server side processing, clustering support, manage huge data (TB), built-in relational capabilities, table-locking or row-locking and can be easily integrated with applications.
- e.g. DB2, Oracle, MS-SQL, MySQL, MS-Access, SQLite, ...
- RDBMS design is based on Codd's rules developed at IBM (in 1970).



Sunbeam Infotech

www.sunbeaminfo.com

SQL

- Clients send SQL queries to RDBMS server and operations are performed accordingly.
- Originally it was named as RQBE (Relational Query By Example).
- SQL is ANSI standardised in 1987 and then revised multiple times adding new features. Recent revision in 2016.
- SQL is case insensitive.
- There are five major categories:
 - DDL: Data Definition Language e.g. CREATE, ALTER, DROP, RENAME.
 - DML: Data Manipulation Language e.g. INSERT, UPDATE, DELETE.
 - DQL: Data Query Language e.g. SELECT.
 - DCL: Data Control Language e.g. CREATE USER, GRANT, REVOKE.
 - TCL: Transaction Control Language e.g. SAVEPOINT, COMMIT, ROLLBACK.
- Table & column names allows alphabets, digits & few special symbols.
- If name contains special symbols then it should be back-quotes.
- e.g. `Tb1`, `T1#`, `T2\$` etc. Names can be max 30 chars long.



Sunbeam Infotech

www.sunbeaminfo.com

MySQL

- Developed by Michael Widenius in 1995. It is named after his daughter name Myia.
- Sun Microsystems acquired MySQL in 2008.
- Oracle acquired Sun Microsystems in 2010.
- MySQL is free and open-source database under GPL. However some enterprise modules are close sourced and available only under commercial version of MySQL.
- MariaDB is completely open-source clone of MySQL.
- MySQL supports multiple database storage and processing engines.
- MySQL versions:
 - < 5.5: MyISAM storage engine
 - 5.5: InnoDB storage engine
 - 5.6: SQL Query optimizer improved, memcached style NoSQL
 - 5.7: Windowing functions, JSON data type added for flexible schema
 - 8.0: CTE, NoSQL document store.
- MySQL is database of year 2019 (in database engine ranking).



Sunbeam Infotech

www.sunbeaminfo.com

MySQL RDBMS

Nilesh Ghule <nilesh@sunbeaminfo.com>

MySQL installation on Ubuntu/Linux

- terminal> sudo apt-get install mysql-community-server mysql-community-client
- This installs MySQL server (mysqld) and MySQL client (mysql).
- MySQL Server (mysqld)
 - Run as background process.
 - Implemented in C/C++.
 - Process SQL queries and generate results.
 - By default run on port 3306.
 - Controlled via systemctl.
 - terminal> sudo systemctl start|stop|status|enable|disable mysql
- MySQL client (mysql)
 - Command line interface
 - Send SQL queries to server and display its results.
 - terminal> mysql -u root -p
- Additional MySQL clients
 - MySQL workbench
 - PHPMyAdmin

Getting started

- root login can be used to perform CRUD as well as admin operations.
- It is recommended to create users for performing non-admin tasks.
 - mysql> CREATE DATABASE db;
 - mysql> SHOW DATABASES;
 - mysql> CREATE USER dbuser@localhost IDENTIFIED BY 'dbpass';
 - mysql> SELECT user, host FROM mysql.user;
 - mysql> GRANT ALL PRIVILEGES ON db.* TO dbuser@localhost;
 - mysql> FLUSH PRIVILEGES;
 - mysql> EXIT;
- terminal> mysql -u dbuser -pdbpass
 - mysql> SHOW DATABASES;
 - mysql> SELECT user(), DATABASE();
 - mysql> USE db;
 - mysql> SHOW TABLES;
 - mysql> CREATE TABLE student(id INT, name VARCHAR(20), marks DOUBLE);
 - mysql> INSERT INTO student VALUES(1, 'Abc', 89.5);
 - mysql> SELECT * FROM student;

Database logical layout

- Database/schema is like a namespace/container that stores all db objects related to a project.
- It contains tables, constraints, relations, stored procedures, functions, triggers, ...
- There are some system databases e.g. mysql, performance_schema, information_schema, sys, ... They contain db internal/system information.
 - e.g. SELECT user, host FROM mysql.user;
- A database contains one or more tables.
- Tables have multiple columns.
- Each column is associated with a data-type.
- Columns may have zero or more constraints.
- The data in table is in multiple rows.
- Each row have multiple values (as per columns).

Database physical layout

- In MySQL, the data is stored on disk in its data directory i.e. /var/lib/mysql
- Each database/schema is a separate sub-directory in data dir.
- Each table in the db, is a file on disk.
- e.g. student table in current db is stored in file /var/lib/mysql/db/student.ibd.
- Data is stored in binary format.
- A file may not be contiguously stored on hard disk.
- Data rows are not contiguous. They are scattered in the hard disk.
- In one row, all fields are consecutive.
- When records are selected, they are selected in any order.

MySQL data types

- RDBMS have similar data types (but not same).
- MySQL data types can be categorised as follows
 - Numeric types (Integers)
 - TINYINT (1 byte), SMALLINT (2 byte), MEDIUMINT (3 byte), INT (4 byte), BIGINT (8 byte), BIT(n bits)
 - integer types can signed (default) or unsigned.
 - Numeric types (Floating point)
 - approx. precision – FLOAT (4 byte), DOUBLE (8 byte) | DECIMAL(m, n) – exact precision
 - Date/Time types
 - DATE, TIME, DATETIME, TIMESTAMP, YEAR
 - String types – size = number of chars * size of char
 - CHAR(1-255) – Fixed length, Very fast access.
 - VARCHAR(1-65535) – Variable length, Stores length + chars.
 - TINYTEXT (255), TEXT (64K), MEDIUMTEXT (16M), LONGTEXT (4G) – Variable length, Slower access.
 - Binary types – size = number of bytes
 - BINARY, VARBINARY, TINYBLOB, BLOB, MEDIUMBLOB, LONGBLOB
 - Miscellaneous types
 - ENUM, SET

CHAR vs VARCHAR vs TEXT

- CHAR
 - Fixed inline storage.
 - If smaller data is given, rest of space is unused.
 - Very fast access.
- VARCHAR
 - Variable inline storage.
 - Stores length and characters.
 - Slower access than CHAR.
- TEXT
 - Variable external storage.
 - Very slow access.
 - Not ideal for indexing.
- CREATE TABLE temp(c1 CHAR(4), c2 VARCHAR(4), c3 TEXT(4));
- DESC temp;
- INSERT INTO temp VALUES('abcd', 'abcd', 'abcdef');

INSERT – DML

- Insert a new row (all columns, fixed order).
 - INSERT INTO table VALUES (v1, v2, v3);
- Insert a new row (specific columns, arbitrary order).
 - INSERT INTO table(c3, c1, c2) VALUES (v3, v1, v2);
 - INSERT INTO table(c1, c2) VALUES (v1, v2);
 - Missing columns data is NULL.
 - NULL is special value and it is not stored in database.
- Insert multiple rows.
 - INSERT INTO table VALUES (av1, av2, av3), (bv1, bv2, bv3), (cv1, cv2, cv3);
- Insert rows from another table.
 - INSERT INTO table SELECT c1, c2, c3 FROM another-table;
 - INSERT INTO table (c1,c2) SELECT c1, c2 FROM another-table;



Sunbeam Infotech

www.sunbeaminfo.com

SQL scripts

- SQL script is multiple SQL queries written into a .sql file.
- SQL scripts are mainly used while database backup and restore operations.
- SQL scripts can be executed from terminal as:
 - terminal> mysql -u user --password db < /path/to/sqlfile
- SQL scripts can be executed from command line as:
 - mysql> SOURCE /path/to/sqlfile
- Note that SOURCE is MySQL CLI client command.
- It reads commands one by one from the script and execute them on server.



Sunbeam Infotech

www.sunbeaminfo.com

SELECT – DQL

- Select all columns (in fixed order).
 - SELECT * FROM table;
- Select specific columns / in arbitrary order.
 - SELECT c1, c2, c3 FROM table;
- Column alias
 - SELECT c1 AS col1, c2 col2 FROM table;
- Computed columns.
 - SELECT c1, c2, c3, expr1, expr2 FROM table;

SELECT c1,
 CASE WHEN condition1 THEN value1,
 WHEN condition2 THEN value2,
 ...
 ELSE valuen
 END
 FROM table;



Sunbeam Infotech

www.sunbeaminfo.com

SELECT – DQL

- Distinct values in column.
 - SELECT DISTINCT c1 FROM table;
 - SELECT DISTINCT c1, c2 FROM table;
- Select limited rows.
 - SELECT * FROM table LIMIT n;
 - SELECT * FROM table LIMIT m, n;




Sunbeam Infotech

www.sunbeaminfo.com

SELECT – DQL – ORDER BY

- In db rows are scattered on disk. Hence may not be fetched in a fixed order.
- Select rows in asc order.
 - SELECT * FROM table ORDER BY c1;
 - SELECT * FROM table ORDER BY c2 ASC;
- Select rows in desc order.
 - SELECT * FROM table ORDER BY c3 DESC;
- Select rows sorted on multiple columns.
 - SELECT * FROM table ORDER BY c1, c2;
 - SELECT * FROM table ORDER BY c1 ASC, c2 DESC;
 - SELECT * FROM table ORDER BY c1 DESC, c2 DESC;
- Select top or bottom n rows.
 - SELECT * FROM table ORDER BY c1 ASC LIMIT n;
 - SELECT * FROM table ORDER BY c1 DESC LIMIT n;
 - SELECT * FROM table ORDER BY c1 ASC LIMIT m, n;



Sunbeam Infotech

www.sunbeaminfo.com

SELECT – DQL – WHERE

- It is always good idea to fetch only required rows (to reduce network traffic).
- The WHERE clause is used to specify the condition, which records to be fetched.
- Relational operators
 - <, >, <=, >=, =, != or <>
- NULL related operators
 - NULL is special value and cannot be compared using relational operators.
 - IS NULL or <>, IS NOT NULL.
- Logical operators
 - AND, OR, NOT



Sunbeam Infotech

www.sunbeaminfo.com

MySQL RDBMS

Nilesh Ghule <nilesh@sunbeaminfo.com>

SELECT – DQL – WHERE

- BETWEEN operator (include both ends)
 - c1 BETWEEN val1 AND val2
- IN operator (equality check with multiple values)
 - c1 IN (val1, val2, val3)
- LIKE operator (similar strings)
 - c1 LIKE 'pattern'.
 - % represent any number of any characters.
 - _ represent any single character.



Sunbeam Infotech

www.sunbeaminfo.com



UPDATE – DML

- To change one or more rows in a table.
- Update row(s) single column.
 - UPDATE table SET c2=new-value WHERE c1=some-value;
- Update multiple columns.
 - UPDATE table SET c2=new-value, c3=new-value WHERE c1=some-value;
- Update all rows single column.
 - UPDATE table SET c2=new-value;



Sunbeam Infotech

www.sunbeaminfo.com



DELETE – DML vs TRUNCATE – DDL vs DROP – DDL

- DELETE
 - To delete one or more rows in a table.
 - Delete row(s)
 - DELETE FROM table WHERE c1=value;
 - Delete all rows
 - DELETE FROM table
- TRUNCATE
 - Delete all rows.
 - TRUNCATE TABLE table;
 - Truncate is faster than DELETE.
- DROP
 - Delete all rows as well as table structure.
 - DROP TABLE table;
 - DROP TABLE table IF EXISTS;
 - Delete database/schema.
 - DROP DATABASE db;

Sunbeam Infotech

www.sunbeaminfo.com



Seeking HELP

- HELP is client command to seek help on commands/functions.
 - HELP SELECT;
 - HELP Functions;
 - HELP SIGN;



Sunbeam Infotech

www.sunbeaminfo.com



DUAL table

- A dummy/in-memory a table having single row & single column.
- It is used for arbitrary calculations, testing functions, etc.
 - SELECT 2 + 3 * 4 FROM DUAL;
 - SELECT NOW() FROM DUAL;
 - SELECT USER(), DATABASE() FROM DUAL;
- In MySQL, DUAL keyword is optional.
 - SELECT 2 + 3 * 4;
 - SELECT NOW();
 - SELECT USER(), DATABASE();

Sunbeam Infotech

www.sunbeaminfo.com



MySQL RDBMS

Nilesh Ghule <nilesh@sunbeaminfo.com>

SQL functions

- RDBMS provides many built-in functions to process the data.
- These functions can be classified as:
 - Single row functions
 - One row input produce one row output.
 - e.g. ABS(), CONCAT(), IFNULL(), ...
 - Multi-row or Group functions
 - Values from multiple rows are aggregated to single value.
 - e.g. SUM(), MIN(), MAX(), ...
- These functions can also be categorized based on data types or usage.
 - Numeric functions
 - String functions
 - Date and Time functions
 - Control flow functions
 - Information functions
 - Miscellaneous functions

Numeric & String functions

- ABS()
- POWER()
- ROUND(), FLOOR(), CEIL()
- ASCII(), CHAR()
- CONCAT()
- SUBSTRING()
- LOWER(), UPPER()
- TRIM(), LTRIM(), RTRIM()
- LPAD(), RPAD()
- REGEXP_LIKE()

Date-Time and Information functions

- VERSION()
- USER(), DATABASE()
- MySQL supports multiple date time related data types
 - DATE (3), TIME (3), DATETIME (5), TIMESTAMP (4), YEAR (1)
- SYSDATE(), NOW()
- DATE(), TIME()
- DAYOFMONTH(), MONTH(), YEAR(), HOUR(), MINUTE(), SECOND(), ...
- DATEDIFF(), DATE_ADD(), TIMEDIFF()
- MAKEDATE(), MAKETIME()

Control and NULL and List functions

- NULL is special value in RDBMS that represents absence of value in that column.
- NULL values do not work with relational operators and need to use special operators.
- Most of functions return NULL if NULL value is passed as one of its argument.
- ISNULL()
- IFNULL()
- NULLIF()
- COALESCE()
- GREATEST(), LEAST()
- IF(condition, true-value, false-value)

Group functions

- Work on group of rows of table.
- Input to function is data from multiple rows & then output is single row. Hence these functions are called as "Multi Row Function" or "Group Functions".
- These functions are used to perform aggregate ops like sum, avg, max, min, count or std dev, etc. Hence these fns are also called as "Aggregate Functions".
- Example: SUM(), AVG(), MAX(), MIN(), COUNT().
- NULL values are ignored by group functions.
- Limitations of GROUP functions:
 - Cannot select group function along with a column.
 - Cannot select group function along with a single row fn.
 - Cannot use group function in WHERE clause/condition.
 - Cannot nest a group function in another group fn.

GROUP BY clause

- GROUP BY is used for analysis of data i.e. generating reports & charts.
- When GROUP BY single column, generated output can be used to plot 2-D chart. When GROUP BY two column, generated output can be used to plot 3-D chart and so on.
- GROUP BY queries are also called as Multi-dimensional / Spatial queries.
- Syntactical Characteristics:
 - If a column is used for GROUP BY, then it may or may not be used in SELECT clause.
 - If a column is in SELECT, it must be in GROUP BY.
- When GROUP BY query is fired on database server, it does following:
 - Load data from server disk into server RAM.
 - Sort data on group by columns.
 - Group similar records by group columns.
 - Perform given aggregate ops on each column.
 - Send result to client.

Transaction

- Transaction is set of DML queries executed as a single unit.
- Transaction examples
 - accounts table [id, type, balance]
 - UPDATE accounts SET balance=balance-1000 WHERE id = 1;
 - UPDATE accounts SET balance=balance+1000 WHERE id = 2;
- RDBMS transaction have ACID properties.
 - Atomicity
 - All queries are executed as a single unit. If any query is failed, other queries are discarded.
 - Consistency
 - When transaction is completed, all clients see the same data.
 - Isolation
 - Multiple transactions (by same or multiple clients) are processed concurrently.
 - Durability
 - When transaction is completed, all data is saved on disk.

Transaction

- Transaction management
 - START TRANSACTION;
 - ...
 - COMMIT WORK;
 - START TRANSACTION;
 - ...
 - ROLLBACK WORK;
- In MySQL autocommit variable is by default 1. So each DML command is auto-committed into database.
 - SELECT @@autocommit;
- Changing autocommit to 0, will create new transaction immediately after current transaction is completed. This setting can be made permanent in config file.
 - SET autocommit=0;

Transaction

- Save-point is state of database tables (data) at the moment (within a transaction).
- It is advised to create save-points at end of each logical section of work.
- Database user may choose to rollback to any of the save-point.
- Transaction management with Save-points
 - START TRANSACTION;
 - ...
 - SAVEPOINT sa1;
 - ...
 - SAVEPOINT sa2;
 - ...
 - ROLLBACK TO sa1;
 - ...
 - COMMIT; // or ROLLBACK
- Commit always commit the whole transaction.
- ROLLBACK or COMMIT clears all save-points.

Transaction

- Transaction is set of DML statements.
- If any DDL statement is executed, current transaction is automatically committed.
- Any power failure, system or network failure automatically rollback current state.
- Transactions are isolated from each other and are consistent.

Row locking


- When an user update or delete a row (within a transaction), that row is locked and becomes read-only for other users.
- The other users see old row values, until transaction is committed by first user.
- If other users try to modify or delete such locked row, their transaction processing is blocked until row is unlocked.
- Other users can INSERT into that table. Also they can UPDATE or DELETE other rows.
- The locks are automatically released when COMMIT/ROLLBACK is done by the user.
- This whole process is done automatically in MySQL. It is called as "OPTIMISTIC LOCKING".

Row locking

- Manually locking the row in advanced before issuing UPDATE or DELETE is known as "PESSIMISTIC LOCKING".
- This is done by appending FOR UPDATE to the SELECT query.
- It will lock all selected rows, until transaction is committed or rolled back.
- If these rows are already locked by another users, the SELECT operation is blocked until rows lock is released.
- By default MySQL does table locking. Row locking is possible only when table is indexed on the column.

Entity Relations

- To avoid redundancy of the data, data should be organized into multiple tables so that tables are related to each other.
- The relations can be one of the following
 - One to One
 - One to Many
 - Many to One
 - Many to Many
- Entity relations is outcome of Normalization process.




Sunbeam Infotech

www.sunbeaminfo.com

SQL Joins

- Join statements are used to SELECT data from multiple tables using single query.
- Typical RDBMS supports following types of joins:
 - Cross Join
 - Inner Join
 - Left Outer Join
 - Right Outer Join
 - Full Outer Join
 - Self join



Sunbeam Infotech

www.sunbeaminfo.com

Cross Join

deptno	dname
10	DEV
20	QA
30	OPS
40	ACC

empno	ename	deptno
1	Amit	10
2	Rahul	10
3	Nilesh	20
4	Nitin	50
5	Sarang	50

- Compares each row of Table1 with every row of Table2.
- Yields all possible combinations of Table1 and Table2.
- In MySQL, The larger table is referred as "Driving Table", while smaller table is referred as "Driven Table". Each row of Driving table is combined with every row of Driven table.
- Cross join is the fastest join, because there is no condition check involved.



Sunbeam Infotech

www.sunbeaminfo.com

Inner Join

deptno	dname
10	DEV
20	QA
30	OPS
40	ACC

empno	ename	deptno
1	Amit	10
2	Rahul	10
3	Nilesh	20
4	Nitin	50
5	Sarang	50

- The inner JOIN is used to return rows from both tables that satisfy the join condition.
- Non-matching rows from both tables are skipped.
- If join condition contains equality check, it is referred as equi-join; otherwise it is non-equi-join.



Sunbeam Infotech

www.sunbeaminfo.com

Left Outer Join

deptno	dname
10	DEV
20	QA
30	OPS
40	ACC

empno	ename	deptno
1	Amit	10
2	Rahul	10
3	Nilesh	20
4	Nitin	50
5	Sarang	50

- Left outer join is used to return matching rows from both tables along with additional rows in left table.
- Corresponding to additional rows in left table, right table values are taken as NULL.
- OUTER keyword is optional.



Sunbeam Infotech

www.sunbeaminfo.com

Right Outer Join

deptno	dname
10	DEV
20	QA
30	OPS
40	ACC

empno	ename	deptno
1	Amit	10
2	Rahul	10
3	Nilesh	20
4	Nitin	50
5	Sarang	50

- Right outer join is used to return matching rows from both tables along with additional rows in right table.
- Corresponding to additional rows in right table, left table values are taken as NULL.
- OUTER keyword is optional.



Sunbeam Infotech

www.sunbeaminfo.com

Full Outer Join

deptno	dname
10	DEV
20	QA
30	OPS
40	ACC

empno	ename	deptno
1	Amit	10
2	Rahul	10
3	Nilesh	20
4	Nitin	50
5	Sarang	50

- Full join is used to return matching rows from both tables along with additional rows in both tables.
- Corresponding to additional rows in left or right table, opposite table values are taken as NULL.
- Full outer join is not supported in MySQL, but can be simulated using set operators.

Set operators

ename	dname
Amit	DEV
Rahul	DEV
Nilesh	QA
NULL	OPS
NULL	ACC

ename	dname
Amit	DEV
Rahul	DEV
Nilesh	QA
Nitin	NULL
Sarang	NULL

- UNION operator is used to combine results of two queries. The common data is taken only once. It can be used to simulate full outer join.
- UNION ALL operator is used to combine results of two queries. Common data is repeated.


Self Join

- When join is done on same table, then it is known as "Self Join". The both columns in condition belong to the same table.
- Self join may be an inner join or outer join.

empno	ename	deptno	mgr
1	Amit	10	4
2	Rahul	10	3
3	Nilesh	20	4
4	Nitin	50	5
5	Sarang	50	NULL

empno	ename	deptno	mgr
1	Amit	10	4
2	Rahul	10	3
3	Nilesh	20	4
4	Nitin	50	5
5	Sarang	50	NULL

Multi-Table Joins



Sub queries

- Sub-query is query within query. Typically it work with SELECT statements.
- Output of inner query is used as input to outer query.
- If no optimization is enabled, for each row of outer query result, sub-query is executed once. This reduce performance of sub-query.
- Single row sub-query
 - Sub-query returns single row.
 - Usually it is compared in outer query using relational operators.

Sub queries

- Multi-row sub-query
 - Sub-query returns multiple rows.
 - Usually it is compared in outer query using operators like IN, ANY or ALL.
 - IN operator compare for equality with results from sub-queries (at least one result should match).
 - ANY operator compares with the results from sub-queries (at least one result should match).
 - ALL operator compares with the results from sub-queries (all results should match).

Sub queries

- Correlated sub-query
 - If number of results from sub-query are reduced, query performance will increase.
 - This can be done by adding criteria (WHERE clause) in sub-query based on outer query row.
 - Typically correlated sub-query use IN, ALL, ANY and EXISTS operators.



Sunbeam Infotech

www.sunbeaminfo.com

Sub query

- Sub queries with UPDATE and DELETE are not supported in all RDBMS.
- In MySQL, Sub-queries in UPDATE/DELETE is allowed, but sub-query should not SELECT from the same table, on which UPDATE/DELETE operation is in progress.



Sunbeam Infotech

www.sunbeaminfo.com

Views

- RDBMS view represents view (projection) of the data.
- View is based on SELECT statement.
- Typically it is restricted view of the data (limited rows or columns) from one or more tables (joins and/or sub-queries) or summary of the data (grouping).
- Data of view is not stored on server hard-disk; but its SELECT statement is stored in compiled form. It speeds up execution of view.



Sunbeam Infotech

www.sunbeaminfo.com

Views

- Views are of two types: Simple view and Complex view
- Usually if view contains computed columns, group by, joins or sub-queries, then the views are said to be complex. DML operations are not supported on these views.
- DML operations on view affects underlying table.
- View can be created with CHECK OPTION to ensure that DML operations can be performed only on the data visible in that view.



Sunbeam Infotech

www.sunbeaminfo.com

View

- Views can be differentiated with: SHOW FULL TABLES.
 - Views can be dropped with DROP VIEW statement.
 - View can be based on another view.
-
- Applications of views
 - Security: Providing limited access to the data.
 - Hide source code of the table.
 - Simplifies complex queries.



Sunbeam Infotech

www.sunbeaminfo.com

Data Control Language

- Security is built-in feature of any RDBMS. It is implemented in terms of permissions (a.k.a. privileges).
- There are two types of privileges.
- System privileges
 - Privileges for certain commands i.e. CREATE, ALTER, DROP, ...
 - Typically these privileges are given to the database administrator or higher authority user.
- Object privileges
 - RDBMS objects are table, view, stored procedure, function, triggers, ...
 - Can perform operations on the objects i.e. INSERT, UPDATE, DELETE, SELECT, CALL, ...
 - Typically these privileges are given to the database users.



Sunbeam Infotech

www.sunbeaminfo.com

MySQL RDBMS

Nilesh Ghule <nilesh@sunbeaminfo.com>

User Management

- User management is responsibility of admin (root).
- New user can be created using CREATE USER.
 - CREATE USER user@host IDENTIFIED BY 'password';
 - host can be hostname of server, localhost (current system) or '%' for all client systems.
- Permissions for the user can be listed using SHOW GRANTS command.
 - SHOW GRANTS FOR user@host;
- Users can be deleted using DROP USER.
 - DROP USER user@host;
- Change user password.
 - ALTER USER user@host IDENTIFIED BY 'new_password';
 - FLUSH PRIVILEGES;



Sunbeam Infotech

www.sunbeaminfo.com

Data Control Language

- Permissions are given to user using GRANT command.
 - GRANT CREATE ON db.* TO user@host;
 - GRANT CREATE ON *.* TO user1@host, user2@host;
 - GRANT SELECT ON db.table TO user@host;
 - GRANT SELECT, INSERT, UPDATE ON db.table TO user@host;
 - GRANT ALL ON db.* TO user@host;
- By default one user cannot give permissions to other user. This can be enabled using WITH GRANT OPTION.
 - GRANT ALL ON *.* TO user@host WITH GRANT OPTION;
- Permissions assigned to any user can be withdrawn using REVOKE command.
 - REVOKE SELECT, INSERT, UPDATE ON db.table FROM user@host;
- Permissions can be activated by FLUSH PRIVILEGES.
 - System GRANT tables are reloaded by this command. Auto done after GRANT, REVOKE.
 - Command is necessary GRANT tables are modified using DML operations.



Sunbeam Infotech

www.sunbeaminfo.com

Index

- Index enable faster searching in tables by indexed columns.
 - CREATE INDEX idx_name ON table(column);
- One table can have multiple indexes on different columns/order.
- Typically indexes are stored as some data structure (like BTREE or HASH) on disk.
- Indexes are updated during DML operations. So DML operation are slower on indexed tables.



Sunbeam Infotech

www.sunbeaminfo.com

Index

- Index can be ASC or DESC.
 - It cause storage of key values in respective order (MySQL 8.x onwards).
 - ASC/DESC index is used by optimizer on ORDER BY queries.
- There are four types of indexes:
 - Simple index
 - CREATE INDEX idx_name ON table(column [ASC|DESC]);
 - Unique index
 - CREATE UNIQUE INDEX idx_name ON table(column [ASC|DESC]);
 - Doesn't allow duplicate values.
 - Composite index
 - CREATE INDEX idx_name ON table(column1 [ASC|DESC], column2 [ASC|DESC]);
 - Composite index can also be unique. Do not allow duplicate combination of columns.
 - Clustered index
 - PRIMARY index automatically created on Primary key for row lookup.
 - If primary key is not available, hidden index is created on synthetic column.
 - It is maintained in tabular form and its reference is used in other indexes.



Sunbeam Infotech

www.sunbeaminfo.com

Index

- Indexes should be created on shorter (INT, CHAR, ...) columns to save disk space.
- Few RDBMS do not allow indexes on external columns i.e. TEXT, BLOB.
- MySQL support indexing on TEXT/BLOB up to n characters.
 - CREATE TABLE test (blob_col BLOB, ..., INDEX(blob_col(10)));
- To list all indexes on table:
 - SHOW INDEXES FROM table;
- To drop an index:
 - DROP INDEX idx_name ON table;
- When table is dropped, all indexes are automatically dropped.
- Indexes should not be created on the columns not used frequent search, ordering or grouping operations.
- Columns in join operation should be indexed for better performance.



Sunbeam Infotech

www.sunbeaminfo.com

Query performance

- Few RDBMS features ensure better query performance.
 - Index speed up execution of SELECT queries (search operations).
 - Correlated sub-queries execute faster.
- Query performance can be observed using EXPLAIN statement.
 - EXPLAIN FORMAT=JSON SELECT ...;
- EXPLAIN statement shows
 - Query cost (Lower is the cost, faster is the query execution).
 - Execution plan (Algorithm used to execute query e.g. loop, semi-join, materialization, etc).
- Optimizations can be enabled or disabled by optimizer_switch system variable.
 - SELECT @@optimizer_switch;
 - SET @@optimizer_switch='materialization=off';



Sunbeam Infotech

www.sunbeaminfo.com

MySQL RDBMS

Nilesh Ghule <nilesh@sunbeaminfo.com>

Constraints

- Constraints are restrictions imposed on columns.
- There are five constraints
 - NOT NULL
 - UNIQUE
 - PRIMARY KEY
 - FOREIGN KEY
 - CHECK
- Few constraints can be applied at either column level or table level. Few constraints can be applied on both.
- Optionally constraint names can be mentioned while creating the constraint. If not given, it is auto-generated.
- Each DML operation check the constraints before manipulating the values. If any constraint is violated, error is raised.

Sunbeam Infotech

www.sunbeaminfo.com

Constraints

- NOT NULL
 - NULL values are not allowed.
 - Can be applied at column level only.
 - CREATE TABLE table(c1 TYPE NOT NULL, ...);
- UNIQUE
 - Duplicate values are not allowed.
 - NULL values are allowed.
 - Not applicable for TEXT and BLOB.
 - UNIQUE can be applied on one or more columns.
 - Internally creates unique index on the column (fast searching).
 - A table can have one or more unique keys.
 - Can be applied at column level or table level.
 - CREATE TABLE table(c1 TYPE UNIQUE, ...);
 - CREATE TABLE table(c1 TYPE, ..., UNIQUE(c1));
 - CREATE TABLE table(c1 TYPE, ..., CONSTRAINT constraint_name UNIQUE(c1));

Sunbeam Infotech

www.sunbeaminfo.com

Constraints

- PRIMARY KEY
 - Column or set of columns that uniquely identifies a row.
 - Only one primary key is allowed for a table.
 - Primary key column cannot have duplicate or NULL values.
 - Internally index is created on PK column.
 - TEXT/BLOB cannot be primary key.
 - If no obvious choice available for PK, composite or surrogate PK can be created.
 - Creating PK for a table is a good practice.
 - PK can be created at table level or column level.
 - CREATE TABLE(table(c1 TYPE PRIMARY KEY, ...);
 - CREATE TABLE table(c1 TYPE, ..., PRIMARY KEY(c1));
 - CREATE TABLE table(c1 TYPE, ..., CONSTRAINT constraint_name PRIMARY KEY(c1));
 - CREATE TABLE table(c1 TYPE, c2 TYPE, ..., PRIMARY KEY(c1, c2));

Sunbeam Infotech

www.sunbeaminfo.com

Constraints

- FOREIGN KEY
 - Column or set of columns that references a column of some table.
 - If column belongs to the same table, it is "self referencing".
 - Foreign key constraint is specified on child table column.
 - FK can have duplicate values as well as null values.
 - FK constraint is applied on column of child table (not on parent table).
 - Child rows cannot be deleted, until parent rows are deleted.
 - MySQL have ON DELETE CASCADE clause to ensure that child rows are automatically deleted, when parent row is deleted. ON UPDATE CASCADE clause does same for UPDATE operation.
 - By default foreign key checks are enabled. They can be disabled by
 - SET @@foreign_key_checks = 0;
 - FK constraint can be applied on table level as well as column level.
 - CREATE TABLE child(c1 TYPE, ..., FOREIGN KEY (c1) REFERENCES parent(col))

Sunbeam Infotech

www.sunbeaminfo.com

Constraints

- CHECK
 - CHECK is integrity constraint in SQL.
 - CHECK constraint specifies condition on column.
 - Data can be inserted/updated only if condition is true; otherwise error is raised.
 - CHECK constraint can be applied at table level or column level.
 - CREATE TABLE table(c1 TYPE, c2 TYPE CHECK condition1, ..., CHECK condition2);

Sunbeam Infotech

www.sunbeaminfo.com

DDL – ALTER statement

- ALTER statement is used to do modification into table, view, function, procedure, ...
- ALTER TABLE is used to change table structure.
- Add new column(s) into the table.
 - ALTER TABLE table ADD col TYPE;
 - ALTER TABLE table ADD c1 TYPE, c2 TYPE;
- Modify column of the table.
 - ALTER TABLE table MODIFY col NEW_TYPE;
- Rename column.
 - ALTER TABLE CHANGE old_col new_col TYPE;
- Drop a column
 - ALTER TABLE DROP COLUMN col;
- Rename table
 - ALTER TABLE table RENAME TO new_table;

Sunbeam Infotech

www.sunbeaminfo.com

www.sunbeaminfo.com

MySQL RDBMS

Nilesh Ghule <nilesh@sunbeaminfo.com>

Normalization

- Concept of table design: Table, Structure, Data Types, Width, Constraints, Relations.
- Goals:
 - Efficient table structure.
 - Avoid data redundancy i.e. unnecessary duplication of data (to save disk space).
 - Reduce problems of insert, update & delete.
- Done from input perspective.
- Based on user requirements.
- Part of software design phase.
- View entire appn on per transaction basis & then normalize each transaction separately.
- Transaction Examples:
 - Banking, Rail Reservation, Online Shopping.

Normalization

- For given transaction make list of all the fields.
- Strive for atomicity.
- Get general description of all field properties.
- For all practical purposes we can have a single table with all the columns. Give meaningful names to the table.
- Assign datatypes and widths to all columns on the basis of general desc of fields properties.
- Remove computed columns.
- Assign primary key to the table.
- At this stage data is in un-normalized form.
- UNF is starting point of normalization.

Normalization

- UNF suffers from
 - Insert anomaly
 - Update anomaly
 - Delete anomaly

Normalization

- 1. Remove repeating group into a new table.
- 2. Key elements will be PK of new table.
- 3. (Optional) Add PK of original table to new table to give us Composite PK.
 - Repeat steps 1-3 infinitely -- to remove all repeating groups into new tables.
 - This is **1-NF**. No repeating groups present here. One to Many relationship between two tables.

Normalization

- 4. Only table with composite PK to be examined.
- 5. Those columns that are not dependent on the entire composite PK, they are to be removed into a new table.
- 6. The key elements on which the non-key elements were originally dependent, it is to be added to the new table, and it will be the PK of new table.
 - Repeat steps 4-6 infinitely -- to separate all non-key elements from all tables with composite primary key.
 - This is **2-NF**. Many-to-Many relationship.

Normalization

- 7. Only non-key elements are examined for inter-dependencies.
- 8. Inter-dependent cols that are not directly related to PK, they are to be removed into a new table.
- 9. (a) Key ele will be PK of new table.
- 9. (b) The PK of new table is to be retained in original table for relationship purposes.
 - Repeat steps 7-9 infinitely to examine all non-key ele from all tables and separate them into new table if not dependent on PK.
 - This is **3-NF**.

Normalization

- To ensure data consistency (no wrong data entered by end user).
- Separate table to be created of well-known data. So that min data will be entered by the end user.
- This is BCNF or 4-NF.



Sunbeam Infotech

www.sunbeaminfo.com

De-normalization

- Normalization will yield a structure that is non-redundant.
- Having too many inter-related tables will lead to complex and inefficient queries.
- To ensure better performance of analytical queries, few rules of normalization can be compromised.
- This process is de-normalization.



Sunbeam Infotech

www.sunbeaminfo.com

Codd's rules

- Proposed by Edgar F. Codd – pioneer of the RDBMS – in 1980.
- If any DBMS follow these rules, it can be considered as RDBMS.
- The 0th rule is the main rule known as "The foundation rule".
 - For any system that is advertised as, or claimed to be, a relational data base management system, that system must be able to manage data bases entirely through its relational capabilities.
- The rest of rules can be considered as elaboration of this foundation rule.



Sunbeam Infotech

www.sunbeaminfo.com

Codd's rules

- Rule 1: The information rule:
 - All information in a relational data base is represented explicitly at the logical level and in exactly one way – by values in tables.
- Rule 2: The guaranteed access rule:
 - Each and every datum (atomic value) in a relational data base is guaranteed to be logically accessible by resorting to a combination of table name, primary key value and column name.
- Rule 3: Systematic treatment of null values:
 - Null values (distinct from the empty character string or a string of blank characters and distinct from zero or any other number) are supported in fully relational DBMS for representing missing information and inapplicable information in a systematic way, independent of data type.



Sunbeam Infotech

www.sunbeaminfo.com

Codd's rules

- Rule 4: Dynamic online catalog based on the relational model:
 - The data base description is represented at the logical level in the same way as ordinary data, so that authorized users can apply the same relational language to its interrogation as they apply to the regular data.
- Rule 5: The comprehensive data sublanguage rule:
 - A relational system may support several languages. However, there must be at least one language that supports all functionalities of a RDBMS i.e. data definition, data manipulation, integrity constraints, transaction management, authorization.



Sunbeam Infotech

www.sunbeaminfo.com

Codd's rules

- Rule 6: The view updating rule:
 - All views that are theoretically updatable are also updatable by the system.
- Rule 7: Possible for high-level insert, update, and delete:
 - The capability of handling a base relation or a derived relation as a single operand applies not only to the retrieval of data but also to the insertion, update and deletion of data.
- Rule 8: Physical data independence:
 - Application programs & terminal activities remain logically unbroken whenever any changes are made in either storage representations or access methods.
- Rule 9: Logical data independence:
 - Application programs & terminal activities remain logically unbroken when information-preserving changes of any kind that theoretically permit un-impaired are made to the base tables.



Sunbeam Infotech

www.sunbeaminfo.com

MySQL RDBMS

Nilesh Ghule <nilesh@sunbeaminfo.com>

Codd's rules

- Rule 10: Integrity independence:
 - Integrity constraints specific to a particular relational database must be definable in the relational data sublanguage and storable in the catalog, not in the application programs.
- Rule 11: Distribution independence:
 - The end-user must not be able to see that the data is distributed over various locations. Users should always get the impression that the data is located at one site only.
- Rule 12: The non-subversion rule:
 - If a relational system has a low-level (single-record-at-a-time) language, that low level cannot be used to subvert or bypass the integrity rules and constraints expressed in the higher level relational language (multiple-records-at-a-time).



Sunbeam Infotech

www.sunbeaminfo.com

MySQL Programming

- RDBMS Programming is an ISO standard – part of SQL standard – since 1992.
- SQL/PSM stands for Persistent Stored Module.
- Inspired from PL/SQL – Programming language of Oracle.
- PSM allows writing programs for RDBMS. The program contains set of SQL statements along with programming constructs e.g. variables, if-else, loops, case, ...
- PSM is a block language. Blocks can be nested into another block.
- MySQL program can be a stored procedure, function or trigger.



Sunbeam Infotech

www.sunbeaminfo.com

MySQL Programming

- MySQL PSM program is written by db user (programmers).
- It is submitted from client, server check syntax & store them into db in compiled form.
- The program can be executed by db user when needed.
- Since programs are stored on server in compiled form, their execution is very fast.
- All these programs will run in server memory.



Sunbeam Infotech

www.sunbeaminfo.com

Stored Procedure

- Stored Procedure is a routine. It contains multiple SQL statements along with programming constructs.
- Procedure doesn't return any value (like void fns in C).
- Procedures can take zero or more parameters.
- Procedures are created using CREATE PROCEDURE and deleted using DROP PROCEDURE.
- Procedures are invoked/called using CALL statement.
- Result of stored procedure can be
 - returned via OUT parameter.
 - inserted into another table.
 - produced using SELECT statement (at end of SP).
- Delimiter should be set before writing procedure.

Sunbeam Infotech

www.sunbeaminfo.com

Stored Procedure

```

CREATE TABLE result(v1 DOUBLE, v2 VARCHAR(50));
-- 01_hello.sql (using editor)
DROP PROCEDURE IF EXISTS sp_hello;
DELIMITER $$

CREATE PROCEDURE sp_hello()
BEGIN
    --
    -- SELECT 1 AS v1, 'Hello World' AS v2;
    END;
$$
DELIMITER ;

CALL sp_hello();
SELECT * FROM result;
    
```



Sunbeam Infotech

www.sunbeaminfo.com

Stored Procedure – PSM Syntax

VARIABLES	IF-ELS	LOOPS	SHOW PROCEDURE
DECLARE varname DATATYPE; DECLARE varname DATATYPE DEFAULT init_value; SET varname = new_value;	IF condition THEN body; END IF; ----- IF condition THEN if-body; ELSE else-body; END IF; ----- IF condition THEN if1-body; ELSE if2-body; ELSE else2-body; END IF; ----- IF condition THEN if1-body; ELSEIF condition THEN if2-body; ELSE else-body; END IF;	WHILE condition DO body; END WHILE; ----- REPEAT body; UNTIL condition END REPEAT; ----- label: IF condition THEN LEAVE label; END IF; ----- CASE WHEN CASE WHEN condition THEN body; WHEN condition THEN body; ELSE body; END CASE;	SHOW PROCEDURE STATUS LIKE '%sp_name%'; SHOW CREATE PROCEDURE sp_name;
CREATE PROCEDURE sp_name(PARAMTYPE p1 DATATYPE) BEGIN ----- -- IN param: Initialized by calling program. -- OUT param: Initialized by called procedure. -- INOUT param: Initialized by calling program and modified by called procedure. -- OUT & INOUT param declared as session variables.	IF condition THEN body; END IF; ----- IF condition THEN if1-body; ELSE if2-body; ELSE else-body; END IF;	DROP PROCEDURE IF EXISTS sp_name;	
CREATE PROCEDURE sp_name(OUT p1 INT) BEGIN SELECT 1 INTO p1; END; SET @res = 0; CALL sp_name(@res); SELECT @res;	----- IF condition THEN if1-body; ELSE if2-body; ELSE else-body; END IF;		

Sunbeam Infotech

www.sunbeaminfo.com

MySQL RDBMS

Nilesh Ghule <nilesh@sunbeaminfo.com>

MySQL Exceptions

- Exceptions are runtime problems, which may arise during execution of stored procedure, function or trigger.
- Required actions should be taken against these errors.
- SP execution may be continued or stopped after handling exception.
- MySQL error handlers are declared as:
 - DECLARE *action* HANDLER FOR *condition* *handler_impl*;
- The *action* can be: CONTINUE or EXIT.
- The *condition* can be:
 - MySQL error code: e.g. 1062 for duplicate entry.
 - SQLSTATE value: e.g. 23000 for duplicate entry, NOTFOUND for end-of-cursor.
 - Named condition: e.g. DECLARE duplicate_entry CONDITION FOR 1062;
- The *handler_impl* can be: Single liner or PSM block i.e. BEGIN ... END;

MySQL Stored Functions

- Stored Functions are MySQL programs like stored procedures.
- Functions can be having zero or more parameters. MySQL allows only IN params.
- Functions must return some value using RETURN statement.
- Function entire code is stored in system table.
- Like procedures, functions allows statements like local variable declarations, if-else, case, loops, etc. One function can invoke another function/procedure and vice-versa. The functions can also be recursive.
- There are two types of functions: DETERMINISTIC and NOT DETERMINISTIC.

CREATE FUNCTION CREATE FUNCTION fn_name(p1 TYPE) RETURNS TYPE [NOT] DETERMINISTIC RECURSIVE body; RETURN value; END;	SHOW FUNCTION SHOW FUNCTION STATUS LIKE 'fn_name'; SHOW CREATE FUNCTION fn_name;
DROP FUNCTION DROP FUNCTION IF EXISTS fn_name;	

MySQL Triggers

- Triggers are supported by all standard RDBMS like Oracle, MySQL, etc.
- Triggers are not supported by WEAK RDBMS like MS-Access.
- Triggers are not called by client's directly, so they don't have args & return value.
- Trigger execution is caused by DML operations on database.
 - BEFORE/AFTER INSERT, BEFORE/AFTER UPDATE, BEFORE/AFTER DELETE.
- Like SP/FN, Triggers may contain SQL statements with programming constructs. They may also call other SP or FN.
- However COMMIT/ROLLBACK is not allowed in triggers. They are executed in same transaction in which DML query is executed.

CREATE TRIGGER CREATE TRIGGER trig_name AFTER BEFORE del_op ON table FOR EACH ROW BEGIN -- body -- use OLD & NEW keywords -- to access old/new rows. -- INSERT triggers - NEW rows. -- DELETE triggers - OLD rows. END;	SHOW TRIGGERS SHOW TRIGGERS FROM db_name;
DROP TRIGGER DROP TRIGGER trig_name;	

MySQL Triggers

- Applications of triggers:
 - Maintain logs of DML operations (Audit Trails).
 - Data cleansing before insert or update data into table. (Modify NEW value).
 - Copying each record AFTER INSERT into another table to maintain "Shadow table".
 - Copying each record AFTER DELETE into another table to maintain "History table".
 - Auto operations of related tables using cascading triggers.
- Cascading triggers
 - One trigger causes execution of 2nd trigger, 2nd trigger causes execution of 3rd trigger and so on.
 - In MySQL, there is no upper limit on number of levels of cascading.
 - This is helpful in complicated business processes.
- Mutating table error
 - If cascading trigger causes one of the earlier trigger to re-execute, "mutating table" error is raised.
 - This prevents infinite loop and also rollback the current transaction.



Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>